

# 1 Towards formally secure compilation of verified F<sup>★</sup> 2 programs against unverified ML contexts 3

4 (Extended Abstract)  
5

6 Cezar-Constantin Andrici<sup>1</sup> Danel Ahman<sup>2</sup> Cătălin Hrițcu<sup>1</sup> Guido Martínez<sup>3</sup>  
7 Abigail Pribisova<sup>1,4</sup> Exequiel Rivas<sup>5</sup> Théo Winterhalter<sup>6</sup>

8 <sup>1</sup>MPI-SP, Bochum, Germany <sup>2</sup>University of Tartu, Estonia <sup>3</sup>Microsoft Research, Redmond, WA, USA  
9 <sup>4</sup>MPI-SWS, Saarbrücken, Germany <sup>5</sup>Tallinn University of Technology, Estonia <sup>6</sup>Inria Saclay, France

10 We are working towards a formally secure compilation framework to compile verified F<sup>★</sup> programs to an ML language.  
11 The framework, itself written and verified in F<sup>★</sup>, aims to compile verified F<sup>★</sup> programs and link them with unverified  
12 ML programs *securely*, so that any unverified code *cannot* inadvertently or maliciously break the internal invariants of  
13 the verified code. In previous work on this project we have built formally secure compilation frameworks [4, 5] between  
14 shallowly embedded subsets of F<sup>★</sup>—i.e., we used shallow embeddings not only for the verified code (which is standard in  
15 F<sup>★</sup>) but also for the unverified code. To compile further to OCaml, we relied on F<sup>★</sup>’s extraction mechanism, which is,  
16 however, unverified. In this extended abstract, we present ongoing work on extending our framework with a formally  
17 secure compilation step to an ML language deeply embedded in F<sup>★</sup>, thus achieving end-to-end security guarantees.

18 This is challenging, since compiling to a deeply embedded  
19 ML language generally involves *quotation*, which is a  
20 meta-program taking a shallowly embedded program and re-  
21 turning a deep embedding of it. If compilation uses quotation,  
22 then end-to-end verification would involve a lot of complex  
23 meta theory, including verifying quotation—proof that was  
24 started for Rocq, but never finished.<sup>1</sup> For our goal of building  
25 an end-to-end secure compiler for F<sup>★</sup> this approach is a non-  
26 starter because there is not enough formalized meta theory  
27 for F<sup>★</sup> and such a formalization would be very challenging,  
28 since F<sup>★</sup> has many verification features that go beyond stan-  
29 dard type theory and also a much larger core than Rocq.

## 42 1 Relational quotation

43 Instead of using quotation, we propose *relational quotation*,  
44 a technique inspired from relational compilation [11, 12] and  
45 work on reflection via canonical structures [8]. Relational  
46 quotation involves a user-defined unary relation on arbitrary  
47 F<sup>★</sup> values that carves out the subset of the host language  
48 that represents the shallow embedding. A derivation of the  
49 relation stands both as a quotation of a program, and *as a*  
50 *proof that it is the quotation of that program*, which eliminates  
51 the necessity of a formalized meta theory.

52 For intuition, here we present a simplified relation C for  
53 Booleans, defined as an inductive relation:

```
54 type C : #a:Type → a → Type =
55 | Cfalse : C false
56 | Ctrue : C true
57 | CIf : #a:Type → #cond:bool → #b1:a → #b2:a →
58   C cond → C b1 → C b2 →
59   C (if cond then b1 else b2)
```

60 To see how this simple relation C works, let us take the  
61 simple program p1 and manually quote it by defining p1C.

```
62 let p1 = if true then false else true
63 let p1C : C p1 = CIf Ctrue Cfalse Ctrue
```

64 With this relation, we see clearly the quotation of p1 (i.e.,  
65 CIf Ctrue Cfalse Ctrue). We can obtain a deep embedding by  
66 recursively defining a projection on the derivation of the  
67 relation.

68 The difference between standard quotation and relational  
69 quotation is that the latter *constructs* a program of type  $\alpha$   
70 that gets unified with the program we are quoting. In other  
71 words, the derivation constructs a program and F<sup>★</sup>’s typing  
72 ensures that it is a correct derivation of the program we  
73 are quoting—thus, eliminating the necessity for additional  
74 meta-theoretic proofs to verify quotation.

75 **Automation.** Above we obtained C p1 manually, but this  
76 was just for intuition. On the one hand, we can write a meta  
77 program that automatically tries to build a derivation. On  
78 the other hand, based on our experiments in F<sup>★</sup>, however,  
79 we did not have to write our own meta program, but **we**  
80 **just defined C as a type class and used type class res-**  
81 **olution!** So we can actually quote program p1 by writing  
82 let p1C : C p1 = solve, where solve is the F<sup>★</sup> tactic that calls  
83 type class resolution. It will be interesting to see if type class  
84 resolution scales with adding more features to the language.

## 85 2 Secure Compilation

86 We define compilation as the meta program that takes a pro-  
87 gram of type  $\alpha$  and returns the deep embedding together with  
88 a proof. We are proving the same compilation criterion as in  
89 our previous work [4, 5], Robust Relational Hyperproperty

90 <sup>1</sup>Verifying quoting in MetaRocq: <https://github.com/MetaRocq/metarocq/blob/coq-8.16/quotation/theories/README.md>

111 Preservation (RrHP), the strongest criterion of Abate et al.  
 112 [1], and also stronger than full abstraction. RrHP ensures  
 113 that compilation preserves observational equivalence, non-  
 114 interference and trace properties. We think such a strong  
 115 criterion should hold for our compiler to an ML language  
 116 after adding more interesting effects (such as state and Input-  
 117 Output) because  $F^*$  itself is designed in the tradition of the  
 118 ML family of languages.

119 Compilation first tries to derive the quotation of the pro-  
 120 gram, this way checking if the program is part of our shallow  
 121 embedding, and if it succeeds, it constructs the deeply em-  
 122 bedded ML program and the proof. The deeply embedded  
 123 program is untyped because the shallow embeddings may  
 124 use dependent types. The proof, however, states that the  
 125 deeply embedded program can be securely linked with arbi-  
 126 trary syntactically typed contexts.

127 We prove RrHP using one more relation, a cross-language  
 128 logical relation between shallowly embedded and deeply em-  
 129 bedded programs. This logical relation is asymmetric, since  
 130 it relates shallowly embedded programs to deeply embedded  
 131 ones, however, its definition is still quite standard though.  
 132 The proof that a shallowly embedded program and its deep  
 133 embedding are in the logical relation is done recursively on  
 134 the derivation of the relation used to do quotation.

135 The current status is that we produce proofs for compiling  
 136 simply typed pure  $F^*$  programs and that are safe to link them  
 137 with syntactically typed pure ML programs.

### 3 Towards end-to-end proofs

141 To be able to connect with our previous work, and be able  
 142 to compile verified effectful programs, we are working on  
 143 the following extensions.

144 *Refinement types.* Refinements are a core feature of  $F^*$   
 145 and they allow one to have functions with pre- and post-  
 146 conditions. The challenge with quoting refined programs is  
 147 that the relation used to do the quotation is bottom-up, it  
 148 constructs a program of type  $\alpha$ , and then it unifies the new  
 149 program with the one we are quoting. This means that when  
 150 we are quoting a function with a pre- and a post-condition,  
 151 the relation constructs the function and it also constructs  
 152 the proof that the post-condition is satisfied. Intuitively, con-  
 153 structing the proof should be guided by the proof done for  
 154 the quoted program, however,  $F^*$  does not save the proofs  
 155 done for refinements, thus, we do not know how to do it (e.g.,  
 156 if one has an  $x$  refined with  $p\ x$ , there is no way to project  
 157 out the proof of  $p\ x$ ). This means that one has to reprove that  
 158 the quoted program satisfies the refinements. Hopefully, this  
 159 can be done by the SMT. Myreen and Owens [11] face the  
 160 same problem when synthesising CakeML programs from  
 161 HOL functions using relational compilation.

162 *Monadic shallow embeddings.* In our previous work, we  
 163 represented verified stateful and IO-performing programs  
 164 using shallow embeddings that use monads. While we did

166 not explore compiling such programs yet, previous work on  
 167 relational compilation shows that such a thing is possible.  
 168 Most notably, the work of Abrahamsson et al. [2] shows how  
 169 to synthesize CakeML programs from HOL functions that  
 170 embed references, exceptions, and IO operations.

171 *Dependent types.* We think this technique could be scaled  
 172 to quotation of dependently typed programs. For that, we  
 173 would probably need a mutually recursive definition between  
 174 a relation for types and one for expressions. A general defi-  
 175 nition would require first-class universe polymorphism (as  
 176 in Agda) which  $F^*$  does not have. Since our goal is to link  
 177 with ML programs, maybe one can limit the definition to a  
 178 limited set of universes.

179 While working on these extensions, one would have to  
 180 also deal with two **challenges** of relational quotation (inher-  
 181 ited from relational compilation) about completeness: (1) one  
 182 is able to relate only what is expressible by a user defined  
 183 relation (e.g., pattern matching is difficult to do in a general  
 184 way), (2) finding the derivation is done using tactics that may  
 185 loop or fail.

## 4 Related work

186 To our knowledge, there is no existing formal secure compiler  
 187 for a proof-oriented languages (e.g., Rocq, Dafny, HOL), but  
 188 there are verified compilers for Rocq and HOL that we discuss  
 189 next.

190 *Verified compilation of Rocq.* Forster et al. [7] present a  
 191 verified compiler for Gallina to Malfunction (an intermediate  
 192 untyped language of OCaml). Their compiler is based on the  
 193 existing MetaRocq [13] project that formalizes Rocq in itself.  
 194 They define compilation as a series of steps, one of them  
 195 being quotation, which is trusted. They prove a realizabil-  
 196 ity relation that guarantees that the extracted untyped code  
 197 operationally behaves as the initial Rocq program. They com-  
 198 pile Rocq without any special step for shallowly embedded  
 199 code using monads, something we would like to do together  
 200 with a proof of security.

201 Similar to Forster et al. [7], there is CertiCoq [3] and Con-  
 202 Cert [6], two compilers for Rocq that target C and blockchain  
 203 languages. Their compilers are partially verified.

204 *Euf* [10] formalizes a subset of Gallina and provides a  
 205 verified compiler to Assembly code. They translate-validate  
 206 quotation to have an end-to-end correctness proof. To be  
 207 able to translate-validate, they require a computational de-  
 208 notation, which restricts what they can quote.

209 *Relational compilation.* Traditionally, compilation is de-  
 210 fined as a function from source programs to target programs.  
 211 Relational compilation treats compilation as a relation, which  
 212 now frees one to search for a target program so that  $\exists t. s \rightarrow t$ .  
 213 The technique we present, relational quotation, is a variant

of relational compilation specialized for the specific task of doing quotation.

Our work is inspired by Pit-Claudel et al. [12] and Myreen and Owens [11] who present how to do relational compilation for Gallina and Isabelle HOL with support for programs shallowly embedded using monads. Their focus is on proving correct compilation, while we focus on proving secure compilation.

Pit-Claudel et al. [12] present a framework for relational compilation designed for performance-critical applications. The framework allows one to build a relational compiler in an extensible and composable way. They show the versatility of the framework by compiling a subset of Gallina to a low-level imperative language.

**CakeML.** Abrahamsson et al. [2] show how starting from a shallowly embedded program in Isabelle HOL one can synthesize an equivalent CakeML program. CakeML [9] is a subset of Standard ML that has a verified compiler that can target ARMv6, ARMv8, x86-x64, MIPS-64, RISC-V and Silver RSA architectures. They have an end-to-end proof of correctness from the original program written in HOL to the low level code that targets a specific architecture. The synthesis they do, however, is a process that happens outside of HOL: it is implemented in Standard ML and it produces a deeply embedded program and a proof of correctness in HOL. In our case, everything happens in F<sup>★</sup>. Even if CakeML does not have a secure compiler, it will be interesting to see if we can target one of the untyped intermediate languages they have to reuse correct compilation.

## References

- [1] Carmine Abate, Roberto Blanco, Deepak Garg, Catalin Hritcu, Marco Patrignani, and Jérémie Thibault. 2019. Journey Beyond Full Abstraction: Exploring Robust Property Preservation for Secure Compilation. In *32nd IEEE Computer Security Foundations Symposium (CSF)*. IEEE, 256–271. doi:10.1109/CSF.2019.00025
- [2] Oskar Abrahamsson, Son Ho, Hrutvik Kanabar, Ramana Kumar, Magnus O. Myreen, Michael Norrish, and Yong Kiam Tan. 2020. Proof-Producing Synthesis of CakeML from Monadic HOL Functions. *Journal of Automated Reasoning (JAR)* (2020). <https://rdcu.be/b4FrU>
- [3] Abhishek Anand, Andrew Appel, Greg Morrisett, Zoe Paraskevopoulou, Randy Pollack, Olivier Savary Belanger, Matthieu Sozeau, and Matthew Weaver. 2017. CertiCoq: A verified compiler for Coq. In *3rd Workshop on Coq for Programming Languages (CoqPL)*. <https://popl17.sigplan.org/details/main/9/CertiCoq-A-verified-compiler-for-Coq>
- [4] Cezar-Constantin Andrici, Danel Ahman, Cătălin Hritcu, Ruxandra Iclanu, Guido Martínez, Exequiel Rivas, and Théo Winterhalter. 2025. SecRef<sup>\*</sup>: Securely Sharing Mutable References between Verified and Unverified Code in F<sup>★</sup>. *Proc. ACM Program. Lang.* 9, ICFP, Article 253 (Aug. 2025), 31 pages. doi:10.1145/3747522
- [5] Cezar-Constantin Andrici, Ştefan Ciobăcă, Catalin Hritcu, Guido Martínez, Exequiel Rivas, Éric Tanter, and Théo Winterhalter. 2024. Securing Verified IO Programs Against Unverified Code in F<sup>★</sup>. *Proc. ACM Program. Lang.* 8, POPL (2024), 2226–2259. doi:10.1145/3632916
- [6] Danil Annenkov, Mikkel Milo, Jakob Botsch Nielsen, and Bas Spitters. 2022. Extracting functional programs from Coq, in Coq.

- [7] Yannick Forster, Matthieu Sozeau, and Nicolas Tabareau. 2024. Verified Extraction from Coq to OCaml. *Proc. ACM Program. Lang.* 8, PLDI (2024), 52–75. doi:10.1145/3656379
- [8] Georges Gonthier, Beta Ziliani, Aleksandar Nanovski, and Derek Dreyer. 2013. How to make ad hoc proof automation less ad hoc. *Journal of Functional Programming* 23, 4 (2013), 357–401.
- [9] Ramana Kumar, Magnus O. Myreen, Michael Norrish, and Scott Owens. 2014. CakeML: a verified implementation of ML. In *The 41st Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL. ACM*, 179–192. doi:10.1145/2535838.2535841
- [10] Eric Mullen, Stuart Pernsteiner, James R. Wilcox, Zachary Tatlock, and Dan Grossman. 2018. Euf: minimizing the Coq extraction TCB. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs* (Los Angeles, CA, USA) (CPP 2018). Association for Computing Machinery, New York, NY, USA, 172–185. doi:10.1145/3167089
- [11] Magnus O. Myreen and Scott Owens. 2014. Proof-producing Translation of Higher-order logic into Pure and Stateful ML. *Journal of Functional Programming (JFP)* 24, 2-3 (May 2014), 284–315. doi:10.1017/S0956796813000282
- [12] Clément Pit-Claudel, Jade Philipoom, Dustin Jamner, Andres Erbsen, and Adam Chlipala. 2022. Relational compilation for performance-critical applications: extensible proof-producing translation of functional models into low-level code. In *PLDI '22: 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation, San Diego, CA, USA, June 13 - 17, 2022*. Ranjit Jhala and Isil Dillig (Eds.). ACM, 918–936. doi:10.1145/3519939.3523706
- [13] Matthieu Sozeau, Yannick Forster, Meven Lennon-Bertrand, Jakob Nielsen, Nicolas Tabareau, and Théo Winterhalter. 2025. Correct and Complete Type Checking and Certified Erasure for Coq, in Coq. *J. ACM* 72, 1, Article 8 (Jan. 2025), 74 pages. doi:10.1145/3706056